

4F-1 MuaLLM: A Multimodal Large Language Model Agent for Circuit Design Assistance with Hybrid Contextual Retrieval-Augmented Generation

Pravallika Abbineni, Saoud Aldowaish, Colin Liechty, Soroosh Noorzad, Morteza Fayazi (University of Utah, Salt Lake City, UT, USA) Ali Ghazizadeh Ghalati (University of Michigan, Ann Arbor, MI, USA)

8A-1 OAH-CIM: Outlier-Aware Hybrid RRAM-SRAM CIM Accelerator with Variation-Robust Sparsity

Zhiwei Zhou, Tong Hu, Han Bao, Houii Zhou, Yuyang Fu, Yi Li, Xianashui Miao (Huazhong University of Science and Technology), Jia Chen, Jiancong Li (Hong Kong University of Science and Technology)

6B-1 SnipSnap: A Joint Compression Format and Dataflow Co-Optimization Framework for Efficient Sparse LLM Accelerator Design

Junyi Wu, Chao Fang, Zhongfeng Wang (Nanjing University)

9A-1 UltraMalloc: Efficient FPGA-based Memory Allocation Framework Optimized for HBM

Yuwei Qu, Yiqing Mao, Yanxiao Jin, Wai-Shing Luk, Lingli Wang (State Key Laboratory of Integrated Chips and Systems, Fudan University, Shanghai, China)

2C-2 Au-MEDAL: Adaptable Grid Router with Metal Edge Detection And Layer Integration

Andrew B. Kahng, Jakana Lee, Doo-seok Yoon (UC San Diego), Seokhveona Kang, Sehyeon Kim (POSTECH)

1C-1 Spiking-NeRF: Neural Graphics Acceleration With Spiking Feature Encoding for Edge 3D Rendering

Jianzhen Gao, Wei Liu, Yue Liu, Hengyi Zhou, Zhiyi Yu, Shanlin Xiao (Sun Yat-sen University, China)

3A-1 DeepPiC: xPU-PIM Cluster Architecture with Adaptive Resource-Aware Task Orchestration for DeepSeek-Style MoE Inference

Zixu Li, Manni Li, Zijian Huang, Jiayu Yang, Wending Zhao, Yinvin Lin, (Fudan University), Chenachen Wang, Haidong Tian, Xiankui Xiong (State Key Laboratory of Mobile Multimedia Technology, ZTE Corporation, Shenzhen, China)

